

# results

March 30, 2025

These experiments were run using two files.

Firstly `continuous.sh` and `train.py`.

Each instance of `train.py` ran a algorithm in a environment one time.

To get 10 seeds per environment with 9 environments and 3 algorithms there are a total of 270 runs that need to be done.

The 270 runs were completed across two job submissions 1-90 and 91-270.

```
[2]: import os
import numpy as np
import pandas as pd
```

## 1 Getting the results

A slight error in the `train.py` code meant that a large chunk of the results had the labels ppo and sac switched, which resulted in the first attempt at the experiments to fail. The result is that the results are split into three different directories from three different grid jobs. Each grid output directory will only be used to get a single model type from.

```
[3]: results_dirs = [f'/home/thompsjame1/grid-output/
↳467649{n}_rl_continuous_control_baseline_experiments' for n in [0,2]]
results_dirs.append('/home/thompsjame1/grid-output/
↳4676867_rl_continuous_control_baseline_experiments')
results_algo = ['td3', 'ppo', 'sac']
results_dirs = [os.path.join(d, "training_evaluations") for d in results_dirs]
results_dirs
```

```
[3]: ['/home/thompsjame1/grid-
output/4676490_rl_continuous_control_baseline_experiments/training_evaluations',
'/home/thompsjame1/grid-
output/4676492_rl_continuous_control_baseline_experiments/training_evaluations',
'/home/thompsjame1/grid-
output/4676867_rl_continuous_control_baseline_experiments/training_evaluations']
```

```
[4]: results_path = [
    os.path.join(results_dir, result_path)
    for results_dir, algo_type in zip(results_dirs, results_algo)]
```

```

    for result_path in os.listdir(results_dir)
    if result_path.startswith(algo_type)
]

len(results_path), results_path[2:5]

```

```

[4]: (270,
      ['/home/thompsjame1/grid-output/4676490_rl_continuous_control_baseline_experiments/training_evaluations/td3_HalfCheetah-v5_1000000_2.npz',
       '/home/thompsjame1/grid-output/4676490_rl_continuous_control_baseline_experiments/training_evaluations/td3_HalfCheetah-v5_1000000_3.npz',
       '/home/thompsjame1/grid-output/4676490_rl_continuous_control_baseline_experiments/training_evaluations/td3_HalfCheetah-v5_1000000_4.npz'])

```

```

[5]: results = []
    for path in results_path:
        if not path.endswith('.npz'):
            continue

        eval_name = path.split('/')[-1]

        model = eval_name.split('_')[0]
        env = eval_name.split('_')[1]
        steps = eval_name.split('_')[2]
        seed = eval_name.split('_')[3].split('.')[0]
        npz = np.load(path)

        results.append({
            'model': model,
            'env': env,
            'total_steps': steps,
            'seed': seed,
            'timesteps': npz['timesteps'],
            'rewards': npz['results'],
            'ep_lengths': npz['ep_lengths'],
        })

    df = pd.DataFrame(results)
    df

```

```

[5]:      model      env total_steps seed \
0      td3  HalfCheetah-v5      1000000      0
1      td3  HalfCheetah-v5      1000000      1
2      td3  HalfCheetah-v5      1000000      2
3      td3  HalfCheetah-v5      1000000      3
4      td3  HalfCheetah-v5      1000000      4
..      ...      ...      ...      ...

```

265	sac	Pusher-v5	1000000	5
266	sac	Pusher-v5	1000000	6
267	sac	Pusher-v5	1000000	7
268	sac	Pusher-v5	1000000	8
269	sac	Pusher-v5	1000000	9

	timesteps \
0	[1000, 2000, 3000, 4000, 5000, 6000, 7000, 800...
1	[1000, 2000, 3000, 4000, 5000, 6000, 7000, 800...
2	[1000, 2000, 3000, 4000, 5000, 6000, 7000, 800...
3	[1000, 2000, 3000, 4000, 5000, 6000, 7000, 800...
4	[1000, 2000, 3000, 4000, 5000, 6000, 7000, 800...
..	...
265	[1000, 2000, 3000, 4000, 5000, 6000, 7000, 800...
266	[1000, 2000, 3000, 4000, 5000, 6000, 7000, 800...
267	[1000, 2000, 3000, 4000, 5000, 6000, 7000, 800...
268	[1000, 2000, 3000, 4000, 5000, 6000, 7000, 800...
269	[1000, 2000, 3000, 4000, 5000, 6000, 7000, 800...

	rewards \
0	[[-2.47477, -1.326504, -0.692302, -1.338596, -...
1	[[-3.266775, -2.67324, -1.939755, -2.038325, -...
2	[[-0.468081, -2.293828, -1.879588, -0.693806, -...
3	[[-1.491101, -2.19, -1.539083, -1.609354, -1.1...
4	[[-2.98534, -2.277705, 0.022336, -1.132101, -2...
..	...
265	[[-47.668813, -55.833692, -67.94624, -51.25912...
266	[[-54.627242, -68.915077, -55.155968, -57.6722...
267	[[-54.29353, -57.748468, -53.698268, -57.73841...
268	[[-54.836784, -52.815107, -59.760734, -58.6271...
269	[[-58.600244, -52.496547, -55.779737, -64.5120...

	ep_lengths
0	[[1000, 1000, 1000, 1000, 1000, 1000, 1000, 10...
1	[[1000, 1000, 1000, 1000, 1000, 1000, 1000, 10...
2	[[1000, 1000, 1000, 1000, 1000, 1000, 1000, 10...
3	[[1000, 1000, 1000, 1000, 1000, 1000, 1000, 10...
4	[[1000, 1000, 1000, 1000, 1000, 1000, 1000, 10...
..	...
265	[[100, 100, 100, 100, 100, 100, 100, 100, 100,...
266	[[100, 100, 100, 100, 100, 100, 100, 100, 100,...
267	[[100, 100, 100, 100, 100, 100, 100, 100, 100,...
268	[[100, 100, 100, 100, 100, 100, 100, 100, 100,...
269	[[100, 100, 100, 100, 100, 100, 100, 100, 100,...

[270 rows x 7 columns]

## 2 Validating the results

I would like to validate two things about the results.

1. That the model types are actually what they say they are.
2. That the correct number of runs and seeds have been run

### 2.1 Validating model type

### 2.2 Checking correct number of runs

## 3 Visualizing the results

```
[8]: df_grouped = df.explode(['timesteps', 'rewards']).groupby(['model', 'env', 'seed', 'timesteps']).agg({
    "rewards": lambda x: np.vstack(x),
}).reset_index()
df_grouped
```

```
[8]:
```

	model	env	timesteps \	
0	ppo	Ant-v5	1000	
1	ppo	Ant-v5	2000	
2	ppo	Ant-v5	3000	
3	ppo	Ant-v5	4000	
4	ppo	Ant-v5	5000	
...	...	...	...	
27004	td3	Walker2d-v5	996000	
27005	td3	Walker2d-v5	997000	
27006	td3	Walker2d-v5	998000	
27007	td3	Walker2d-v5	999000	
27008	td3	Walker2d-v5	1000000	
				rewards
0				[[994.605428, 994.789186, 994.98996, 990.99214...
1				[[997.7473, 984.486736, 995.396178, 995.329621...
2				[[991.231405, 989.554447, 986.973473, 987.5504...
3				[[989.156628, 985.691269, 982.338319, 987.3692...
4				[[978.897593, 979.452877, 989.408207, 987.4948...
...				...
27004				[[3586.390619, 3328.577188, 3469.250313, 3408...
27005				[[3077.539078, 3103.373637, 3082.026325, 3024...
27006				[[3202.509432, 3237.036607, 3191.531424, 3249...
27007				[[3358.502255, 3330.063387, 3323.810543, 3330...
27008				[[3329.715101, 3312.109436, 3322.453729, 3316...

[27009 rows x 4 columns]

```
[9]: df_grouped["mean_rewards"] = df_grouped["rewards"].apply(lambda x: np.mean(x))
df_grouped["max_rewards"] = df_grouped["rewards"].apply(lambda x: np.max(x))
df_grouped["min_rewards"] = df_grouped["rewards"].apply(lambda x: np.min(x))
df_grouped
```

```
[9]:      model      env  timesteps \
0      ppo      Ant-v5      1000
1      ppo      Ant-v5      2000
2      ppo      Ant-v5      3000
3      ppo      Ant-v5      4000
4      ppo      Ant-v5      5000
...
27004  td3  Walker2d-v5    996000
27005  td3  Walker2d-v5    997000
27006  td3  Walker2d-v5    998000
27007  td3  Walker2d-v5    999000
27008  td3  Walker2d-v5   1000000
```

```
                                rewards  mean_rewards \
0      [[994.605428, 994.789186, 994.98996, 990.99214...  993.814065
1      [[997.7473, 984.486736, 995.396178, 995.329621...  994.207132
2      [[991.231405, 989.554447, 986.973473, 987.5504...  987.942770
3      [[989.156628, 985.691269, 982.338319, 987.3692...  988.130807
4      [[978.897593, 979.452877, 989.408207, 987.4948...  984.505891
...
27004  [[3586.390619, 3328.577188, 3469.250313, 3408...  3750.588397
27005  [[3077.539078, 3103.373637, 3082.026325, 3024...  3674.283195
27006  [[3202.509432, 3237.036607, 3191.531424, 3249...  3473.941411
27007  [[3358.502255, 3330.063387, 3323.810543, 3330...  3744.905717
27008  [[3329.715101, 3312.109436, 3322.453729, 3316...  3741.679792
```

```
                                max_rewards  min_rewards
0      1004.952310    985.528747
1      1003.606960    984.486736
2      1006.239413    965.776645
3      1004.213439    958.616273
4      1014.399710    964.516934
...
27004  4787.474235    888.866982
27005  4833.346584    2126.675978
27006  4904.334854    1696.109651
27007  4885.951055    1829.380533
27008  4889.393571    2786.645381
```

```
[27009 rows x 7 columns]
```

```

[10]: models = df['model'].unique()
      envs = df['env'].unique()
      seeds = df['seed'].unique()

[11]: import matplotlib.pyplot as plt

      for env in envs:
          plt.figure(figsize=(10, 6))
          for model in models:
              df_plot = df_grouped[(df_grouped['model'] == model) &
          ↪(df_grouped['env'] == env)]
              if df_plot.empty:
                  continue

              # Plot the mean rewards as a line
              plt.plot(df_plot['timesteps'], df_plot['mean_rewards'], label=model,
          ↪alpha=0.7)

              # Fill between the min and max rewards to create the shaded area
              plt.fill_between(
                  df_plot['timesteps'],
                  df_plot['min_rewards'],
                  df_plot['max_rewards'],
                  alpha=0.2
              )

              plt.title(f'Average Reward for {env}')
              plt.xlabel('Timesteps')
              plt.ylabel('Average Reward')
              plt.legend()
              plt.grid(True)
              plt.show()

```











